

# Functional Prototype Demonstration 2

Team Epsilon-Greedy Quants

Michael Lee, Nikat Patel, Jose Antonio Alatorre Sanchez

## What we did in this milestone?

We implemented the Policy Gradient algorithms REINFORCE with baseline and Actor-Critic. We also refactored our model code to work in PyTorch.

## Presentation Overview

- REINFORCE with Baseline Summary
- Actor-Critic Summary
- Model Performance Overview
- Discussion of Problems Encountered
- Code Documentation/Organization
- Next Steps

# REINFORCE Summary

## Policy Gradient Method

- Estimates Policy directly, not from Action-Value function
- Continuous action space

## REINFORCE

- Performance under Policy-Gradient Theorem:  
$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a (q_\pi(s, a)) \nabla \pi(a|s, \theta)$$
- Relies on estimated return by Monte-Carlo method
- Uses episode samples to update policy parameter  $\theta$
- **High variance results in slow learning**

## REINFORCE with Baseline

- Compares the action-value to an arbitrary baseline  $b(s)$ 
  - Performance under Policy-Gradient Theorem:
$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a (q_\pi(s, a) - b(s)) \nabla \pi(a|s, \theta)$$
  - Can be any function or random variable as long as it does not vary with action  $a$
  - Commonly used baseline: state value function  $\hat{v}(S_t, w)$
  - Policy parameter  $\theta$  is updated using baseline:
$$\theta_{t+1} = \theta_t + \alpha (G_t - b(S_t)) \frac{\nabla \pi(A_t|S_t, \theta_t)}{\pi(A_t|S_t, \theta_t)}$$
- Baseline functions don't change expected value of update but **can reduce the variance** (speed up learning)

## REINFORCE with Baseline Algorithm Steps

From Sutton and Barto, Chapter 13.4

Steps:

- Initialize the policy parameter  $\theta$  and state-value weights  $w$  at random.
- Loop forever (for each episode):
  - Generate one episode using policy  $\pi_\theta : S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ .
  - Loop for each step of the episode  $t=0,1,\dots,T-1$ :
    - Estimate the return  $G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k$
    - **Calculate the delta between  $G$  and baseline function:**  
 $\delta \leftarrow G - \hat{v}(S_t, w)$
    - **Update the state-value weights:**  $w \leftarrow w + \alpha^w \delta \nabla \hat{v}(S_t, w)$
    - Update policy parameters:  $\theta \leftarrow \theta + \alpha^\theta \gamma_t \delta \nabla \ln \pi(A_t | S_t, \theta)$ 
      - $\alpha^\theta$  - stepsize
      - $\gamma$  - discount factor
      - $\nabla \ln \pi(A_t | S_t, \theta)$  - eligibility vector: gradient of the probability of taking action  $A_t$  given a state  $S_t$  and policy  $\pi_\theta$

# Actor-Critic Methods

In REINFORCE with baseline, the learned state-value function estimates the value of the only the first state of each state transition. This estimate sets a baseline for the subsequent return, but is made prior to the transition's action and thus cannot be used to assess that action. In actor-critic methods, on the other hand, the state-value function is applied also to the second state of the transition. The estimated value of the second state, when discounted and added to the reward, constitutes the one-step return,  $G_{t:t+1}$  which is a useful estimate of the actual return and thus is a way of assessing the action.

When the state-value function is used to assess actions in this way it is called a critic, and the overall policy-gradient method is termed an actor-critic method. Note that the bias in the gradient estimate is not due to bootstrapping as such; the actor would be biased even if the critic was learned by a Monte Carlo method.

## One-step Actor-critic

One-step actor-critic methods replace the full return of with the one-step return (and use a learned state-value function as the baseline) as follows

$$\begin{aligned}\boldsymbol{\theta}_{t+1} &= \boldsymbol{\theta}_t + \alpha(G_{t:t+1} - \hat{v}(S_t, \boldsymbol{w})) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta})}{\pi(A_t | S_t, \boldsymbol{\theta})} \\ \boldsymbol{\theta}_{t+1} &= \boldsymbol{\theta}_t + \alpha(R_{t+1} \gamma \hat{v}(S_{t+1}, \boldsymbol{w}) - \hat{v}(S_t, \boldsymbol{w})) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta})}{\pi(A_t | S_t, \boldsymbol{\theta})}\end{aligned}$$

The main appeal of one-step methods is that they are fully online and incremental, yet avoid the complexities of eligibility traces. They are a special case of the eligibility trace methods, but easier to understand

## One step pseudo code:

### One-step Actor–Critic (episodic), for estimating $\pi_{\theta} \approx \pi_*$

Input: a differentiable policy parameterization  $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization  $\hat{v}(s, \mathbf{w})$

Parameters: step sizes  $\alpha^{\theta} > 0$ ,  $\alpha^{\mathbf{w}} > 0$

Initialize policy parameter  $\theta \in \mathbb{R}^{d'}$  and state-value weights  $\mathbf{w} \in \mathbb{R}^d$  (e.g., to  $\mathbf{0}$ )

Loop forever (for each episode):

    Initialize  $S$  (first state of episode)

$I \leftarrow 1$

    Loop while  $S$  is not terminal (for each time step):

$A \sim \pi(\cdot|S, \theta)$

        Take action  $A$ , observe  $S', R$

$\delta \leftarrow R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$       (if  $S'$  is terminal, then  $\hat{v}(S', \mathbf{w}) \doteq 0$ )

$\mathbf{w} \leftarrow \mathbf{w} + \alpha^{\mathbf{w}} \delta \nabla \hat{v}(S, \mathbf{w})$

$\theta \leftarrow \theta + \alpha^{\theta} I \delta \nabla \ln \pi(A|S, \theta)$

$I \leftarrow \gamma I$

$S \leftarrow S'$

# Eligibility Traces Actor- Critic

The generalizations to the forward view of step methods and then to a  $\lambda$ -return algorithm are straightforward. The one-step return in (1) is merely replaced by  $G_{t:t+1}$  or  $G_t^\lambda$  respectively. The backward view of the  $\lambda$ -return algorithm is also straightforward, using separate eligibility traces for the actor and critic. Pseudocode for the complete algorithm is given in the box below

## Actor-Critic with Eligibility Traces (episodic), for estimating $\pi_\theta \approx \pi_*$

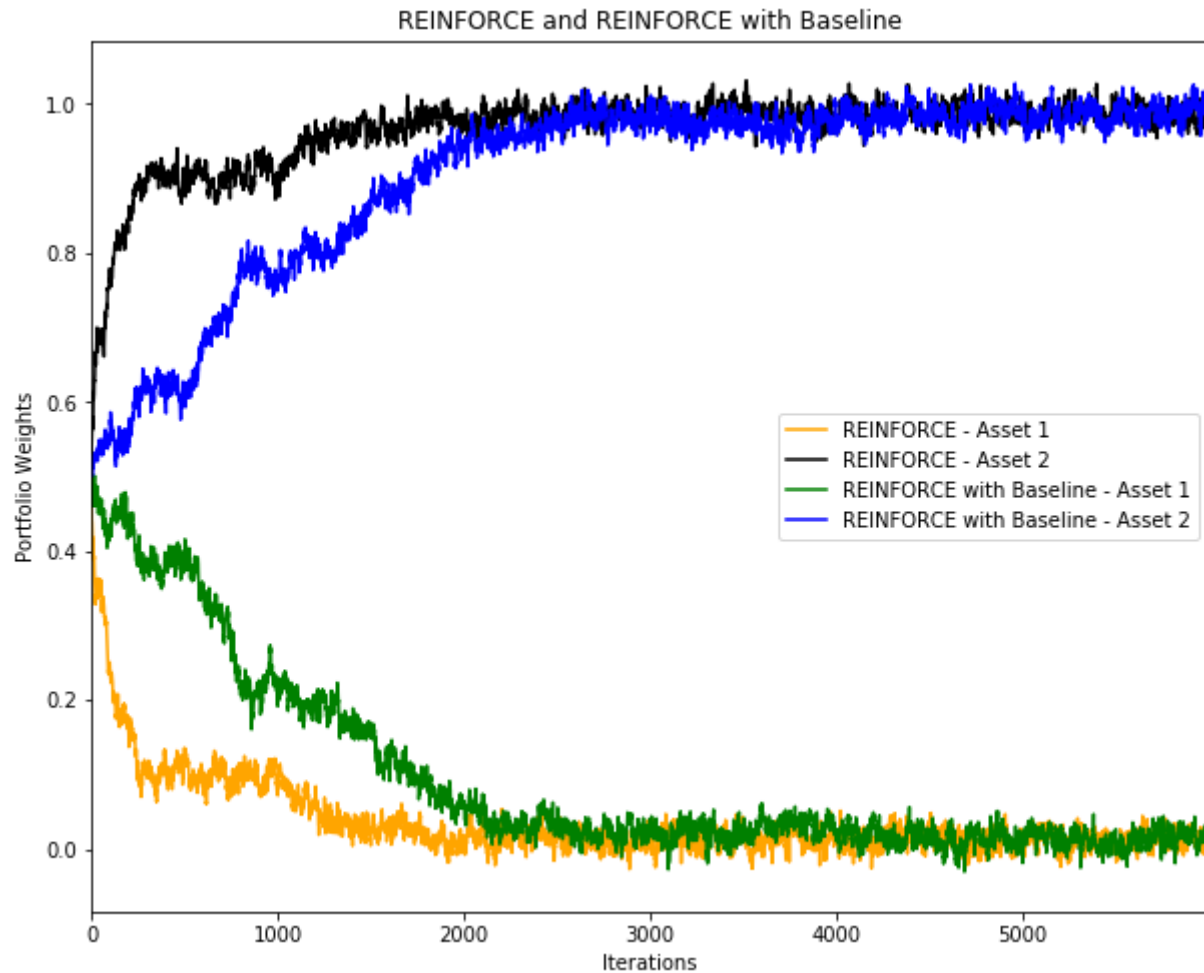
```
Input: a differentiable policy parameterization  $\pi(a|s, \theta)$ 
Input: a differentiable state-value function parameterization  $\hat{v}(s, \mathbf{w})$ 
Parameters: trace-decay rates  $\lambda^\theta \in [0, 1]$ ,  $\lambda^\mathbf{w} \in [0, 1]$ ; step sizes  $\alpha^\theta > 0$ ,  $\alpha^\mathbf{w} > 0$ 
Initialize policy parameter  $\theta \in \mathbb{R}^{d'}$  and state-value weights  $\mathbf{w} \in \mathbb{R}^d$  (e.g., to  $\mathbf{0}$ )
Loop forever (for each episode):
  Initialize  $S$  (first state of episode)
   $\mathbf{z}^\theta \leftarrow \mathbf{0}$  ( $d'$ -component eligibility trace vector)
   $\mathbf{z}^\mathbf{w} \leftarrow \mathbf{0}$  ( $d$ -component eligibility trace vector)
   $I \leftarrow 1$ 
  Loop while  $S$  is not terminal (for each time step):
     $A \sim \pi(\cdot|S, \theta)$ 
    Take action  $A$ , observe  $S', R$ 
     $\delta \leftarrow R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})$  (if  $S'$  is terminal, then  $\hat{v}(S', \mathbf{w}) \doteq 0$ )
     $\mathbf{z}^\mathbf{w} \leftarrow \gamma \lambda^\mathbf{w} \mathbf{z}^\mathbf{w} + \nabla \hat{v}(S, \mathbf{w})$ 
     $\mathbf{z}^\theta \leftarrow \gamma \lambda^\theta \mathbf{z}^\theta + I \nabla \ln \pi(A|S, \theta)$ 
     $\mathbf{w} \leftarrow \mathbf{w} + \alpha^\mathbf{w} \delta \mathbf{z}^\mathbf{w}$ 
     $\theta \leftarrow \theta + \alpha^\theta \delta \mathbf{z}^\theta$ 
     $I \leftarrow \gamma I$ 
     $S \leftarrow S'$ 
```



# REINFORCE Results

REINFORCE - { Orange, Black }


REINFORCE with Baseline - { Green, Blue }



# Actor Critic Results

Actor Critic Without Eligibility Traces - { Orange, Black }


Actor Critic With Eligibility Traces - { Green, Blue }



# Actor Critic Results

Actor Critic Without Eligibility Traces - { Orange, Black }

Actor Critic With Eligibility Traces - { Green, Blue }



## Problems Encountered with TensorFlow

- Issue: TensorFlow REINFORCE model would not converge
  - Used tensorflow.GradientTape() for automatic differentiation
  - Experimented with various keras optimizers
- Solution: Refactored Environment to support PyTorch
  - PyTorch REINFORCE model converged much faster than our standard REINFORCE model

## Version Control Repository

- We have created a private GitHub repository that contains all our data, documentations, code, and notebooks.
- We use version control to develop, update, and collaborate our work.

```
— README.md
— data
— data_env
  — ief.parquet
  — spy.parquet
— lib
  — Benchmarks.py
  — DataHandling.py
  — Environment.py
  — Environment_refactored.py
  — __init__.py
— notebooks
  — Benchmark_EDA.ipynb
  — ENVIRONMENT.ipynb
  — ACTOR_CRITIC.ipynb
  — REINFORCE_BASELINE.ipynb
  — REINFORCE.ipynb
  — data_handler_test.ipynb
  — environment_gymai.ipynb
  — sharpe_sample.ipynb
— static
```

## Next Steps

- Experiment with Various Reward Functions to observe differences
  - Sortino Ratio to Control Drawdowns. Increase weight of negative rewards if drawdown reaches a certain threshold.
- Begin Testing our Models using Real-World DataSets.
- Begin Capstone Documentation.